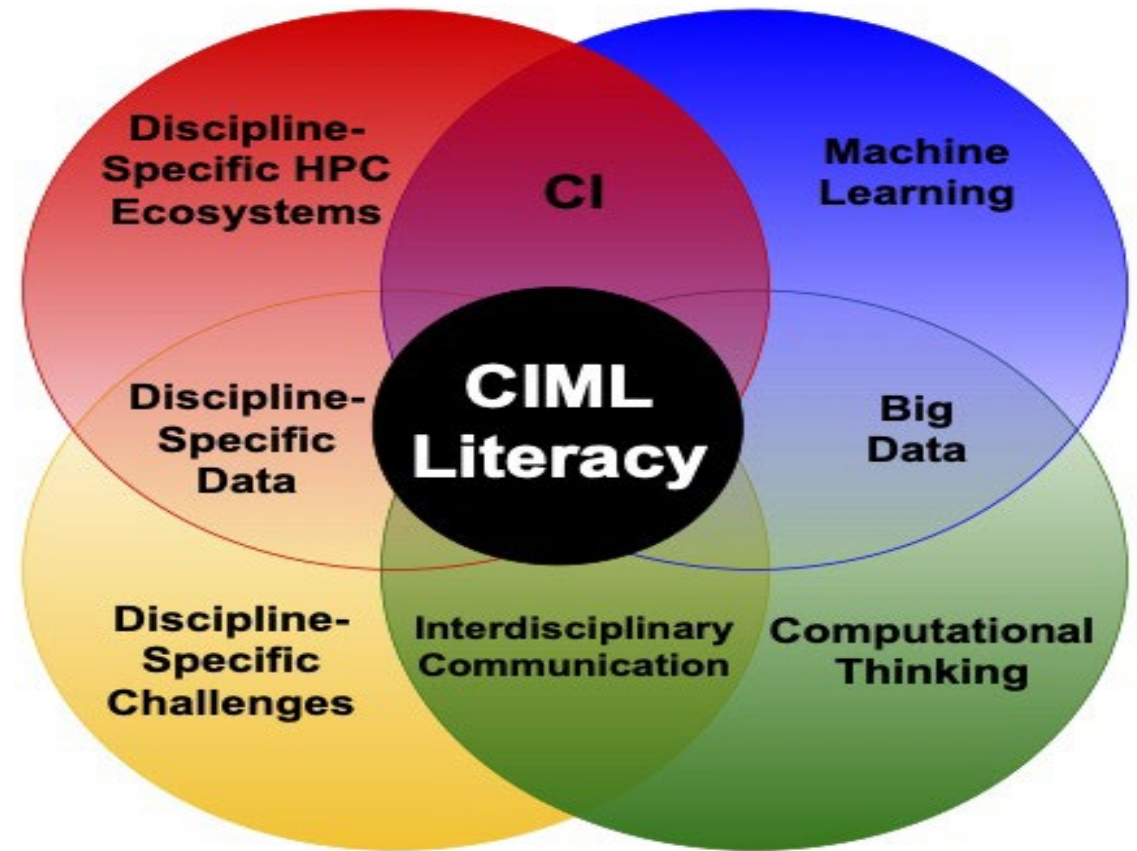


# Scalable Machine Learning: R (and other packages), HPC, and Scaling

*July 2022*

**Paul Rodriguez, PhD  
(SDSC)**



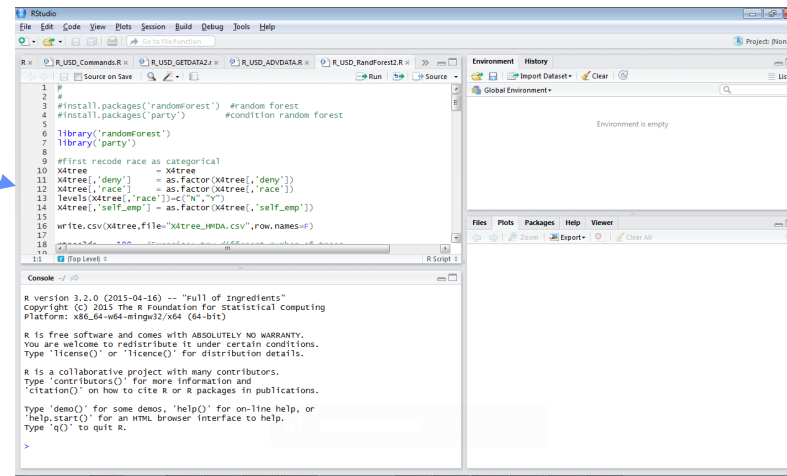
# Outline

- **R and Scaling**
- **Parallel R**
- **Embarassingly Parallel R**
- **A big data exploration of R**

# A typical R development workflow

- R studio: An Integrated development environment for R

Menu tab  
Edit window to  
Build scripts  
R console



Environment  
Information on  
variables and  
command history

Plots, help  
docs, package  
lists

# Typical R code workflow

## #READ DATA

```
X = read.csv('hmda_aer.csv', header=T, stringsAsFactors=T)
```

## #SUBSET DATA

```
indices_2keep = which(X[, 's13'] %in% c(3,4,5))  
X = X[unique(indices_2keep),]
```

## #CREATE/TRANSFORM VARIABLES

```
pi_rat = as.numeric(X[, 's46']/100) #debt2income ratio
```

## #RUN MODEL and SHOW RESULTS

```
lm_result = lm(deny~pi_rat) #lm is 'linear model'  
summary(lm_result)
```

# R strengths for HPC (IMHO)

- **Data Wrangling –**
- **Particular statistical procedure implementations -**
  - Imputation methods (for missing data)
  - Sampling methods
  - Instrument Variable (2 stage) Regression
  - Matching subjects for pairwise analysis
  - MCMC routines (but Stan is likely better package)
  - Generalized Linear Model
  - Some ML model (e.g. randomForest)

# R Scaling In a nutshell

- R uses BLAS/LAPACK math libraries for operations on vectors  
*[Same for Matlab and Python]*
- R packages provide multicore, out-of-core, multinode, or distributed data (SparkR) options  
*[Same for Matlab and Python]*
- Some ML model implementations may be built to use parallel backends (review the available options)

# R Scaling In a nutshell

- R sessionInfo() command shows math libraries


```
> sessionInfo()
R version 4.0.2 (2020-06-22)
Platform: x86_64-pc-linux-gnu (64-bit)
Running under: CentOS Linux 8 (Core)

Matrix products: default
BLAS/LAPACK: /cm/shared/apps/spack/cpu/opt/spack/linux-centos8-zen2/gcc-9.2.0/openblas
grrcfayp3br6kmcnelbgrepqmadwv43e/lib/libopenblas_zenp-r0.3.10.so

locale:
 [1] LC_CTYPE=en_US.UTF-8      LC_NUMERIC=C
 [3] LC_TIME=en_US.UTF-8      LC_COLLATE=en_US.UTF-8
 [5] LC_MONETARY=en_US.UTF-8  LC_MESSAGES=en_US.UTF-8
```

# Consider Regression Computations

- **Linear Model:**  $Y = X * B$   
where  $Y$ =outcomes ,  $X$ =data matrix
- **Algebraically, we could:**  
take “inverse” of  $X * Y = B$  (time consuming)  
use derivatives to search for solutions (very general)
- **Or, better:**  
QR decomposition of  $X$  into triangular matrices (easier to solve but more memory)





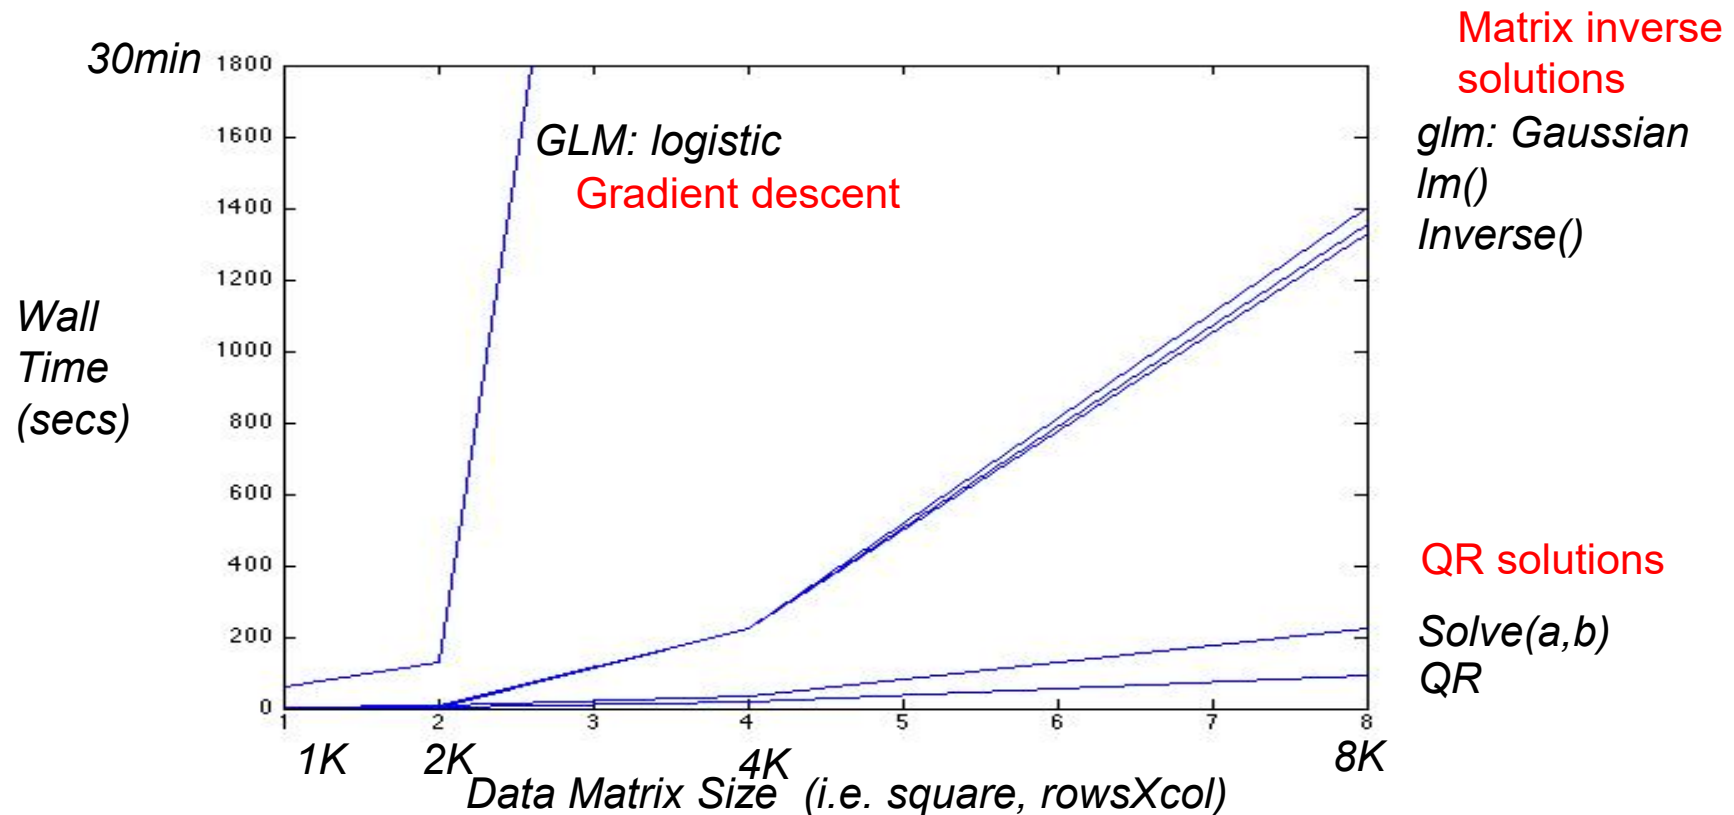
# Solving Linear Systems

## Performance with R, 1 compute node

R:

`glm(Y~X,family=gaussian)` #gaussn regrssn (like `lm`)

`glm(Y~X,family=binomial)` # logistic regrssn (Y=0 or 1)



# Outline

- R and Scaling
- **Parallel R**
- Embarassingly Parallel R
- A big data exploration of R

# R multicore processing


- ‘doParallel’ package – provides the back end to the ‘for each’ parallel processing command
- uses threads across cpu cores to pass data & commands
- It also works for multinode (runs on top of RMPI)

*See <https://cran.r-project.org/web/packages/doParallel/vignettes/gettingstartedParallel.pdf>*

# R multicore coding

```
install.packages(doParallel)  
library(doParallel)  
registerDoParallel(cores=24)
```

**1. allocate workers**



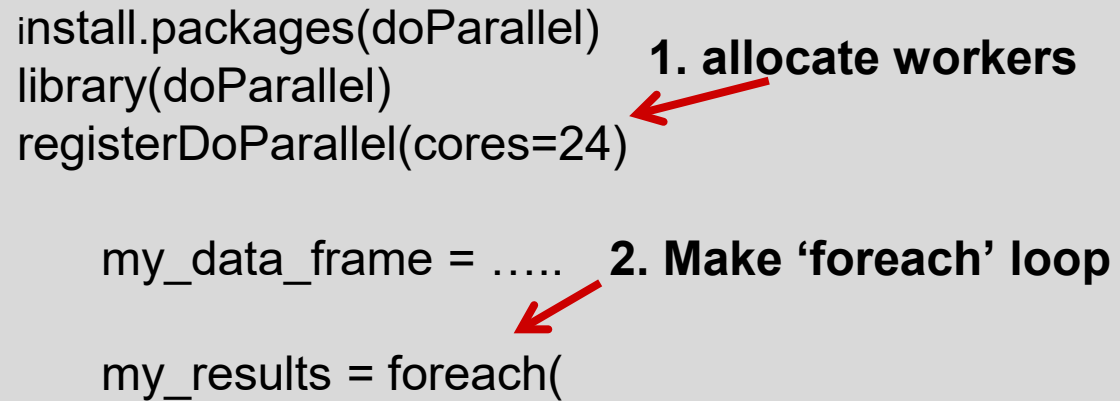
# R multicore coding

```
install.packages(doParallel)
library(doParallel)
registerDoParallel(cores=24)
```

**1. allocate workers**

**2. Make 'foreach' loop**

```
my_data_frame = .....
my_results = foreach(
```



# R multicore coding

```
install.packages(doParallel)  
library(doParallel)  
registerDoParallel(cores=24)
```


**1. allocate workers**



```
my_data_frame = ..... 2. Make 'foreach' loop
```

```
my_results = foreach(i=1:24,.combine=rbind)
```

**3. specify how to  
combine results**



# R multicore coding

```
install.packages(doParallel)
library(doParallel)
registerDoParallel(cores=24)

my_data_frame = .....
my_results = foreach(i=1:24,.combine=rbind) %dopar%
{ ... }
```

**1. allocate workers**

**2. Make 'foreach' loop**

**3. specify how to combine results**

**4. %dopar% runs it across cores, (%do% runs it serially)**

# R multicore coding

```
install.packages(doParallel)
library(doParallel)
registerDoParallel(cores=24)

my_data_frame = .....
my_results = foreach(i=1:24,.combine=rbind) %dopar%
{
  ...
  your code here

  return( a variable or object )
}
```

**1. allocate workers**

**2. Make 'foreach' loop**

**3. specify how to combine results**

**4. %dopar% runs it across cores, (%do% runs it serially)**



# R multicore coding

```
install.packages(doParallel)
library(doParallel)
registerDoParallel(cores=24)
```

**1. allocate workers**

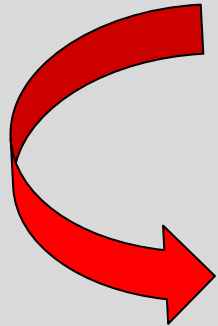
```
my_data_frame = .....
```

**2. Make 'foreach' loop**

```
my_results = foreach(i=1:24,.combine=rbind) %dopar%
{
  ...
  your code here
  return( a variable or object )
}
```

**3. specify how to combine results**

**4. %dopar% runs it across cores, (%do% runs it serially)**



**BEWARE:** foreach will copy data to every core if its seems necessary

# R multinode: parallel backend

```
library(doParallel)
```

```
cl <- makeCluster(48)  
registerDoParallel(cl)
```

**1. allocate cluster as  
parallel backend**



# R multinode: parallel backend

```
library(doParallel)
```

```
cl <- makeCluster(48)  
registerDoParallel(cl)
```


**1. allocate cluster as  
parallel backend**



```
my_data_frame = .....
```

```
results = foreach(i=1:48,.combine=rbind) %dopar%  
{ ... your code here
```

**2.  
%dopar% puts  
loops across  
cores and  
nodes**



```
    return( a variable or object )  
  })  
stopCluster(cl)
```

# R multinode: parallel backend

```
library(doParallel)
```

```
cl <- makeCluster(48)  
registerDoParallel(cl)
```


**1. allocate cluster as  
parallel backend**



```
my_data_frame = .....
```

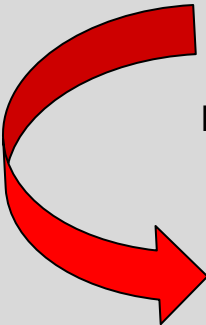
```
results = foreach(i=1:48,.combine=rbind) %dopar%  
{ ... your code here
```

**2.  
%dopar% puts  
loops across  
cores and  
nodes**



```
return( a variable or object )
```

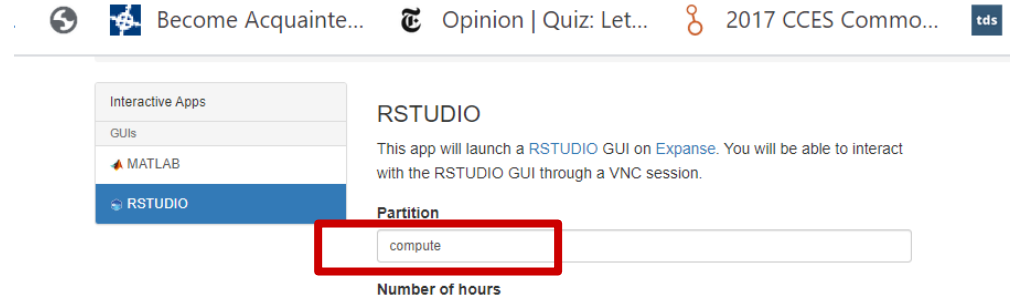
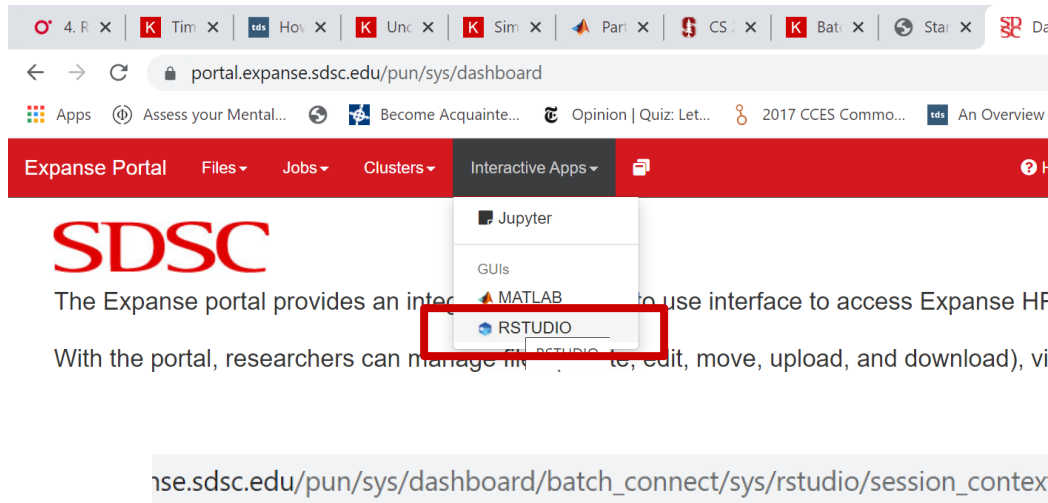
```
})  
stopCluster(cl)
```



**BEWARE: foreach will copy data to every core in every node if its seems necessary**

# Testing/Evaluating R parallel

- **Exercise: 'TestdoParallel' R script**
  1. **Log into expanse portal and start R studio**
  2. **Also log into expanse command line and ssh to compute node**
  3. **run 'top -u username' to see performance**
    - Enter 'H' to see threads; Enter f, toggle 'P' to see cpu ids
    - look for tradeoffs in memory vs execution as matrix size varies (see next slides)



```

-----
Last login: Fri Jun  4 15:01:29 2021 from 71.128.8.73
[p4rodrig@login02 ~]$ squeue -u p4rodrig
              JOBID PARTITION    NAME             USER  ST       TIME  NODES NODELIST(REASON)
              3246260    compute Sys/dash p4rodrig    R       0:27      1 exp-2-15
[p4rodrig@login02 ~]$ ssh exp-2-15
Last login: Sat Jun  5 13:09:04 2021
[p4rodrig@exp-2-15 ~]$ top -u $USER

```

1  
Open portal ->  
Interactive Apps ->  
Rstudio

Enter  
Node: "compute"  
Cores: "64"  
Memory: 124 Gb  
(other fields defaults ok)

2  
Also login:  
login.exppanse

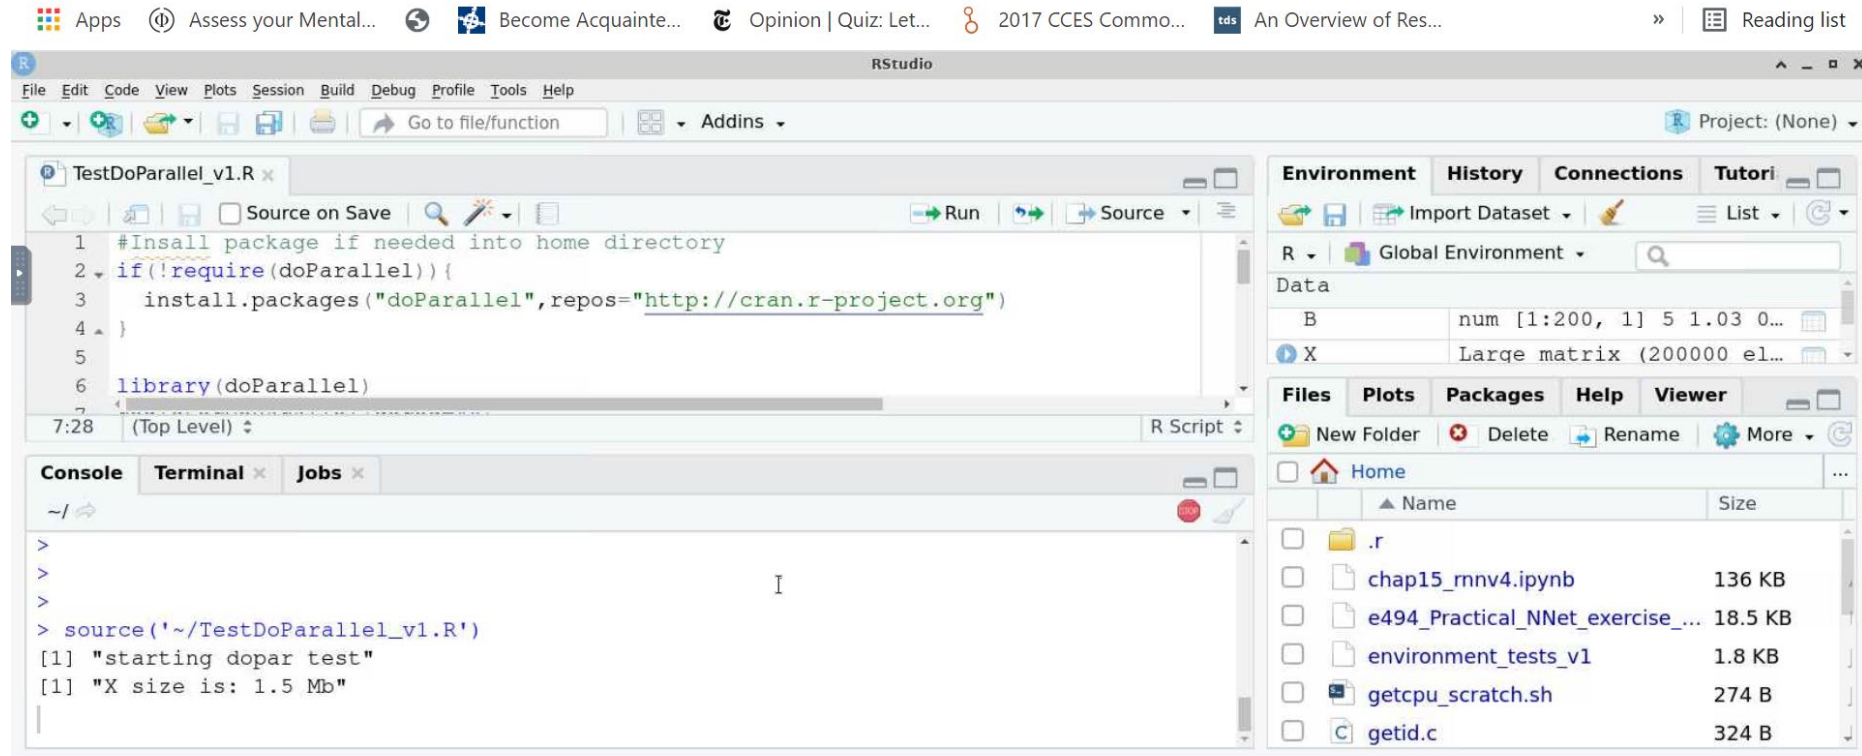
`$ squeue -u $USER`  
`$ ssh exp-##-##`  
`$ top -u $USER`

'H' will toggle threads

'f', highlight P, space, esc  
will toggle last cpuid

Open the 'Test\_doParallel' Rscript

Select 'source' to run the whole script



The screenshot displays the RStudio environment with the following components:

- Top Bar:** Includes browser tabs for 'Apps', 'Assess your Mental...', 'Become Acquainte...', 'Opinion | Quiz: Let...', '2017 CCES Commo...', 'tds An Overview of Res...', and a 'Reading list' icon.
- Menu Bar:** Standard RStudio menus: File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, Help.
- Toolbar:** Icons for file operations and a 'Go to file/function' search bar.
- Source Editor:** Contains the script 'TestDoParallel\_v1.R' with the following code:

```
1 #Install package if needed into home directory
2 if(!require(doParallel)){
3   install.packages("doParallel", repos="http://cran.r-project.org")
4 }
5
6 library(doParallel)
```

The script is set to 'Source on Save' and the 'Source' button is highlighted.
- Console:** Shows the output of the script execution:

```
>
>
>
> source('~/.TestDoParallel_v1.R')
[1] "starting dopar test"
[1] "X size is: 1.5 Mb"
```
- Environment Panel:** Shows the 'Global Environment' with variables 'B' (a numeric vector) and 'X' (a large matrix).
- Files Panel:** Displays a file explorer view of the home directory, listing files such as '.r', 'chap15\_rnnv4.ipynb', 'e494\_Practical\_NNet\_exercise...', 'environment\_tests\_v1', 'getcpu\_scratch.sh', and 'getid.c'.

Review the top output

```
p4rodrig@exp-9-27:~$ top - 15:35:51 up 87 days, 21:07, 1 user, load average: 20.97, 5.76, 4.37
Tasks: 1788 total, 41 running, 1747 sleeping, 0 stopped, 0 zombie
%Cpu(s): 32.6 us, 0.1 sy, 0.0 ni, 67.0 id, 0.0 wa, 0.2 hi, 0.1 si, 0.0 st
MiB Mem : 257517.8 total, 210968.9 free, 42132.2 used, 4416.7 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used, 210846.5 avail Mem

  PID USER      PR  NI  VIRT  RES  SHR S %CPU  %MEM    TIME+  COMMAND
 54587 p4rodrig  20   0 17.5g 896016 4196 R 100.0  0.3   0:37.27 rsession
 54562 p4rodrig  20   0 17.5g 896016 4196 R  99.7  0.3   0:37.48 rsession
 54568 p4rodrig  20   0 17.5g 896016 4196 R  99.7  0.3   0:37.46 rsession
 54571 p4rodrig  20   0 17.5g 896016 4196 R  99.7  0.3   0:37.35 rsession
 54572 p4rodrig  20   0 17.5g 896016 4196 R  99.7  0.3   0:37.33 rsession
 54574 p4rodrig  20   0 17.5g 896016 4196 R  99.7  0.3   0:37.38 rsession
 54579 p4rodrig  20   0 17.5g 896016 4196 R  99.7  0.3   0:37.33 rsession
 54591 p4rodrig  20   0 17.5g 896016 4196 R  99.7  0.3   0:37.24 rsession
```

Notice the elapsed time and memory size

Change the  $N \times P$  matrix size and rerun

(start with  $N=10K$ ,  $P=2K$ )

```
7:28 (Top Level)
Console Terminal x Jobs x
~/
>
> source('~/.TestDoParallel_v1.R')
[1] "starting dopar test"
[1] "X size is: 1.5 Mb"
      user system elapsed
1.176   1.969  30.620
>
```

```
10 # Make up some random data and lis
11 N=100000;      #N rows start with
12 P=2000;        #P columns 200 for
13
14 #make random data with 1 column ar
15 X =matrix(rnorm(N*P),N,P)
16 X[,1] =X[,1]+1
17
16:28 (Top Level)
Console Terminal x Jobs x
~/
Loading required package: foreach
Loading required package: iterators
Loading required package: parallel
[1] "starting dopar test"
[1] "X size is: 1.5 Gb"
```



*Try this at home:*

*Let  $N=100K$ ,  $P=2000$*

*Notice the memory used is close to 124Gb we asked for*

```
p4rodrig@exp-9-27:~  
top - 15:38:40 up 87 days, 21:10, 1 user, load average: 10.77, 6.29, 4.76  
Tasks: 1749 total, 19 running, 1730 sleeping, 0 stopped, 0 zombie  
%Cpu(s): 14.0 us, 0.0 sy, 0.0 ni, 85.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st  
MiB Mem : 257517.8 total, 130239.0 free, 123199.7 used, 4079.0 buff/cache  
MiB Swap: 0.0 total, 0.0 free, 0.0 used, 129947.3 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND	P
55219	p4rodrig	20	0	24.2g	7.6g	2696	R	100.0	3.0	0:24.52	rsession	68
55227	p4rodrig	20	0	24.2g	7.6g	3064	R	100.0	3.0	0:24.55	rsession	88
55235	p4rodrig	20	0	24.2g	7.6g	2696	R	100.0	3.0	0:24.56	rsession	80
55236	p4rodrig	20	0	24.2g	7.6g	2696	R	100.0	3.0	0:24.70	rsession	100
55237	p4rodrig	20	0	24.2g	7.6g	2696	R	100.0	3.0	0:24.50	rsession	47
55242	p4rodrig	20	0	24.2g	7.6g	2696	R	100.0	3.0	0:24.36	rsession	32
55253	p4rodrig	20	0	24.2g	7.6g	2696	R	100.0	3.0	0:24.69	rsession	126
55259	p4rodrig	20	0	24.2g	7.6g	2696	R	100.0	3.0	0:24.00	rsession	16
55261	p4rodrig	20	0	24.2g	7.6g	2696	R	100.0	3.0	0:24.25	rsession	24
55265	p4rodrig	20	0	24.2g	7.6g	2696	R	100.0	3.0	0:23.96	rsession	6
55239	p4rodrig	20	0	24.2g	7.6g	2696	R	99.7	3.0	0:24.61	rsession	20
55241	p4rodrig	20	0	24.2g	7.6g	2696	R	99.7	3.0	0:24.43	rsession	8
55243	p4rodrig	20	0	24.2g	7.6g	2836	R	99.7	3.0	0:24.53	rsession	104

*If you ask for 248Gb will it run?*

*What if you use only 24 cores?*

# Parallezing for loops

(pseudo code)

## R with doParallel

*makecluster*  
*registercluster*

*foreach with dopar,*

*combine results*

## Matlab with parallel toolbox

*parcluster('local')*  
*parpool()*

*parfor*  
*or*  
*'spmd' with*  
*distributed arrays*

*gather array*

## Python with dask.distributed

Import delayed, Client  
Client(numwkr)

for i in range(numwkr):  
  A=delayed(my\_func)(i)  
  Acombine.append(A)


Acombined.compute()

# Outline

- R and Scaling
- Parallel R
- **Embarassingly Parallel R**
- A big data exploration of R

# An option for (embarrassingly) Parallel R


1. Split up data into N parts



# An option for (embarrassingly) Parallel R

1. Split up  
data into N  
parts

2. In slurm batch script:  
`mpirun ... my-perl-script`



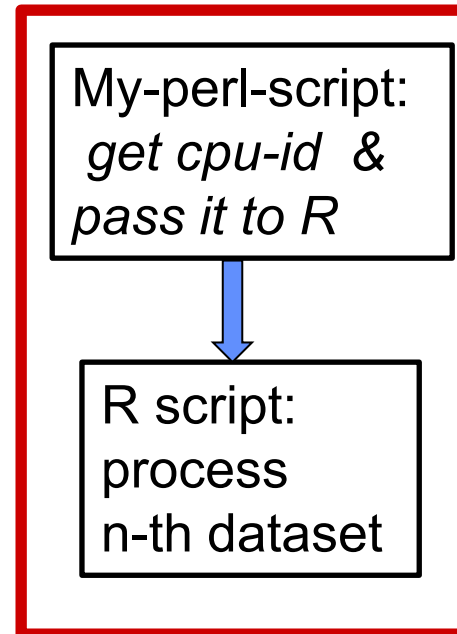

My-perl-script:  
*get cpu-id &  
pass it to R*

R script:  
process  
n-th dataset

# An option for (embarrassingly) Parallel R

1. Split up  
data into N  
parts

2. In slurm batch script:  
`mpirun ... my-perl-script`



*This gets  
distributed  
across nodes  
and cores by  
slurm & mpi  
parameters*

# Slurm parameters: one R instance per core across all nodes

Normal  
batch  
job info

```
...  
#SBATCH --partition=compute  
#SBATCH --nodes=2  
#SBATCH --ntasks-per-node=128  
#SBATCH --cpus-per-task=1
```

2 x 128 = 256 mpi ranks

```
module load slurm  
module load cpu  
module load gcc  
module load intel-mpi
```

256 perl script/R instances  
1 core each


```
mpirun -genv I_MPI_PIN_DOMAIN=omp:compact ./get_mpirank_runcmd.pl
```

(based on /cm/shared/examples/sdsc/mpi-openmp-hybrid/hybrid-slurm.sb)

# one R instance per core across all nodes

1. Split up  
data into N  
parts

2. In slurm batch script:  
`mpirun ... my-perl-script`

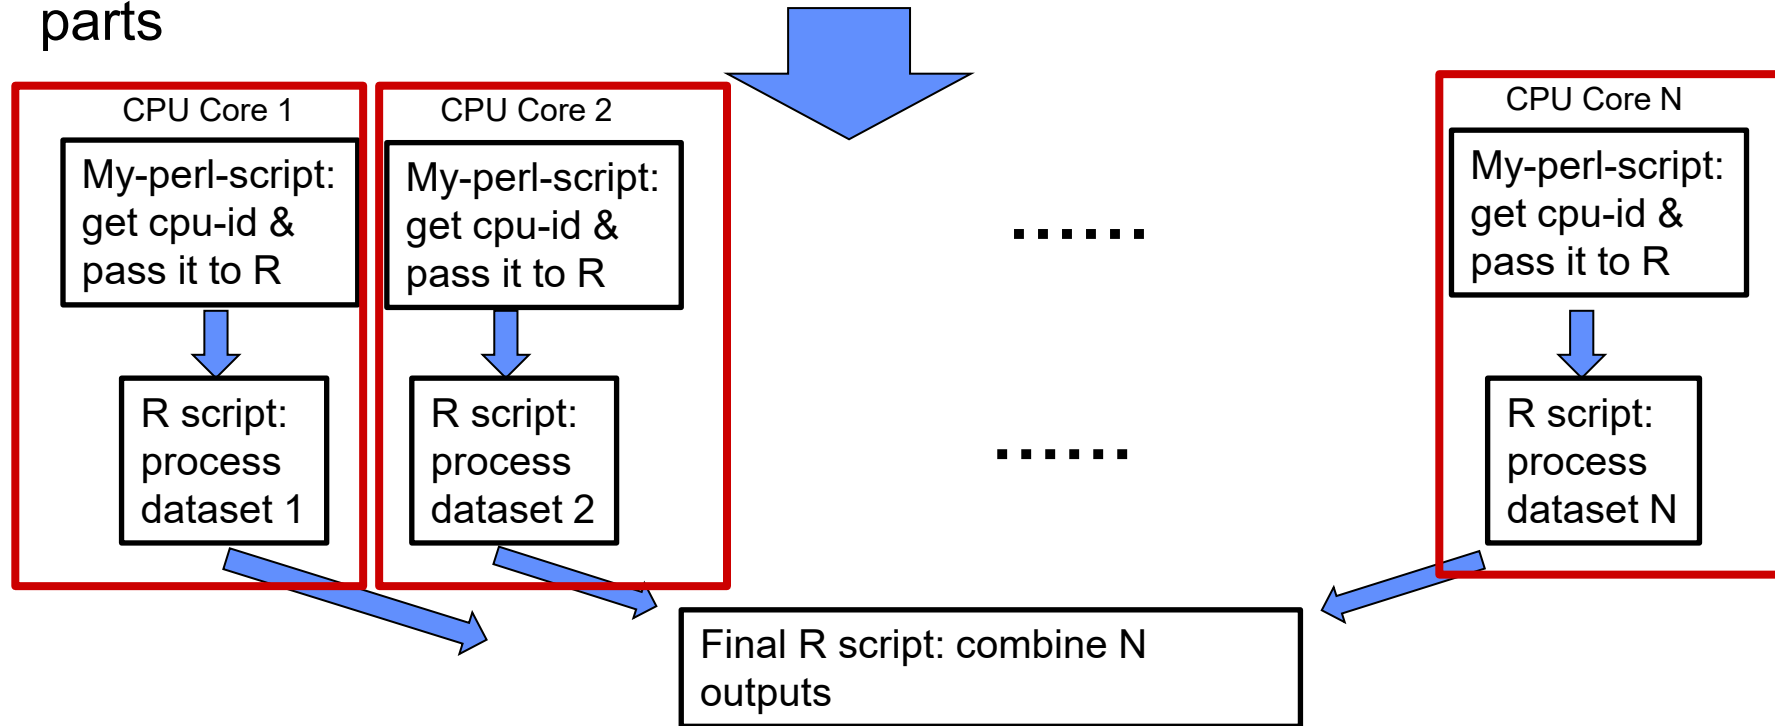




# one R instance per core across all nodes

1. Split up  
data into N  
parts

2. In slurm batch script:  
mpirun ... my-perl-script



*More programming but perhaps more useful*

# Slurm parameters: one R instance per node with 128 cores per R instance

Normal  
batch  
job info

```
...  
#SBATCH --partition=compute  
#SBATCH --nodes=2  
#SBATCH --ntasks-per-node=1  
#SBATCH --cpus-per-task=128  
  
module load slurm  
module load cpu  
module load gcc  
module load intel-mpi  
  
module load r  
  
mpirun -genv I_MPI_PIN_DOMAIN=omp:compact ./get_mpirank_runcmd.pl
```

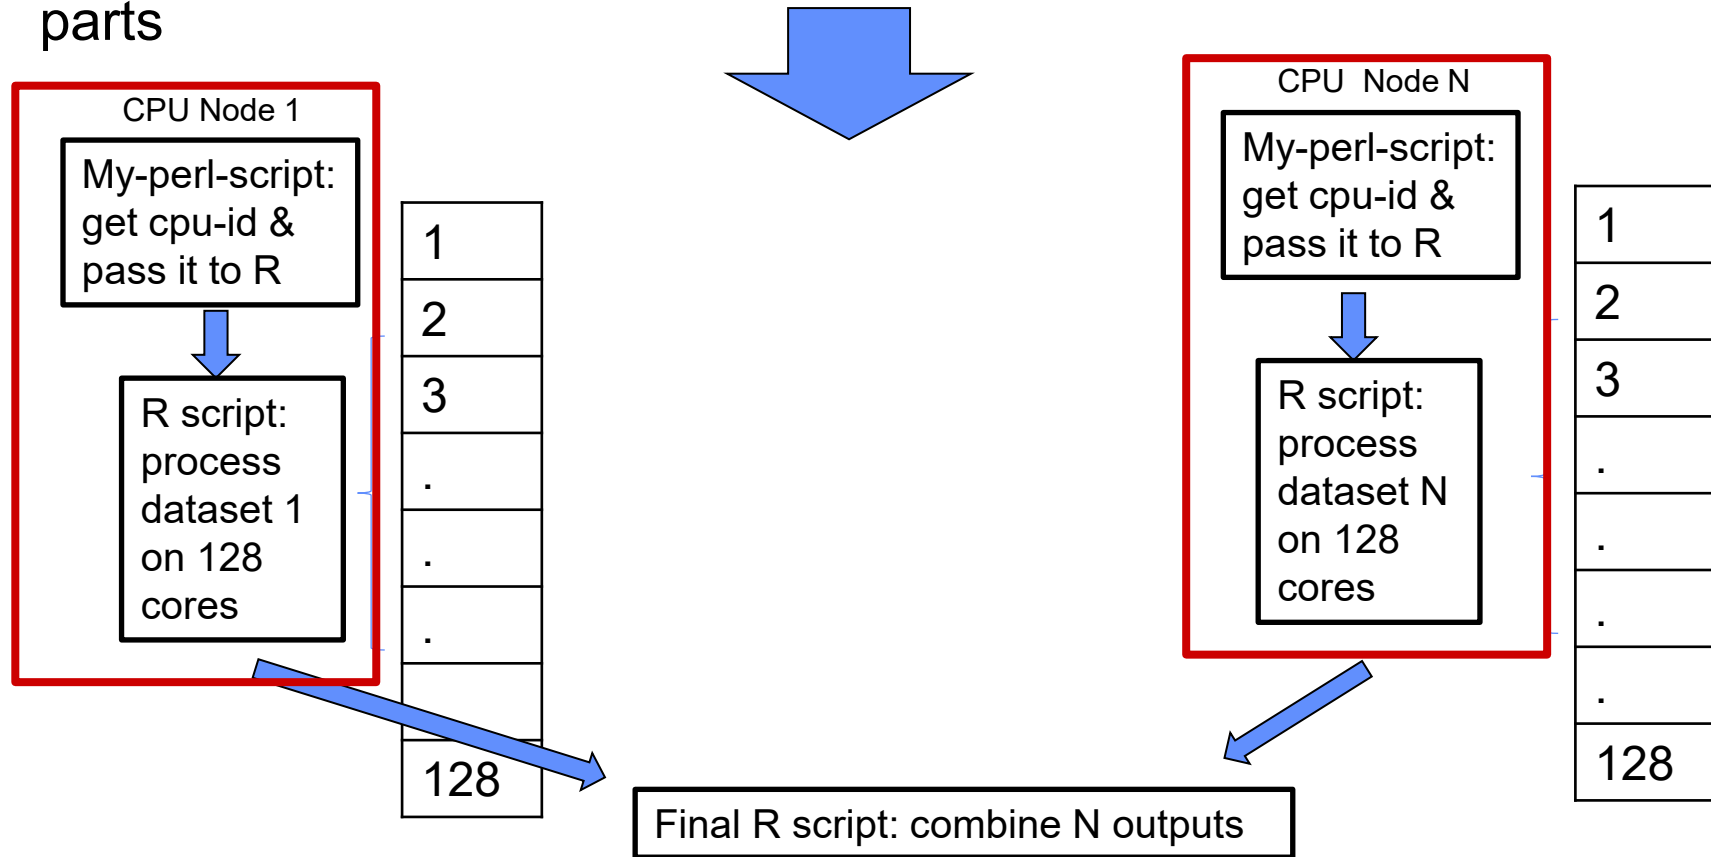
2 x 1 = 2 mpi ranks

2 perl script/R instances  
128 cores each  
(doParallel can use them)

# Example: One R instance per node, doParallel across all cores in each node

1. Split up  
data into N  
parts

2. In slurm batch script:  
mpirun ... my-perl-script



# Outline

- R and Scaling
- Parallel R
- Embarassingly Parallel R
- **A big data exploration of R**

# Big Data exploration

- Run R 'biglasso' with a dataset too big for RAM
- Create large CSV file (117Gb) of X data matrix and Y outcomes:  
 $Y = X * B + \text{noise}$  (where X is 100K x 50K)
- Explore other packages with out-of-core dataset functions:  
Matlab, Dask-ML, Spark, Keras
- Beware: lasso implementations can differ;  
also, Y, X, B should be 'nice'

# Lasso Regression

- **Penalized Loss function**  $L = MSE + \lambda \sum |b_i|$
- **Recall: using penalty term is the same as using a constraint (constrained optimization)**

*find min  $MSE$  such that  $\sum |b_i| < S$*

- **Different implementations may:**
  - use different fit methods (ie forward stepwise, coordinate descent, gradient descent,...)
  - might parallelize by splitting up data, computations, or vectorizing
  - read/load data more or less efficiently

# Considerations

- All packages have special functions to handle out of core dataset
- Sometimes better to have Y,X together; sometimes two files are better
- All were run as Expanse batch jobs, not through notebook/portal
- All were set up run on 1 compute node (248Gb RAM) and use scratch (ssd) space. (Large memory nodes are available on Expanse, but not tested)
- Mostly default parameters used; Little optimization performed; Not a benchmark study!

# R

- R – biglasso (bigmatrix) package to set up file backed dataframe  
<https://cran.rstudio.com/web/packages/biglasso/>

**Issue: the file backend for the big data object path option was hard to get right – ended up just running out of scratch SSD as working directory;**

**copies everything into binary file and descriptor file and got results in about 2hours**



# R code highlights

Use biglasso package

Y data fits in memory so just read it in

X data will be setup with file-backed memory

biglasso() arguments look like the glmnet implementation

```
library(biglasso)
```

```
....
```

```
Y.bm=read.big.matrix(inputYfile,sep = ",")
```

```
X.bm=setupX(inputXfile,sep=","
```

```
type="double",
```

```
backingfile = "x.bin",
```

```
descriptorfile = "x.desc")
```

```
....
```

```
bl_results=biglasso( X.bm, Y.bm,
```

```
row.idx = 1:nrow(X.bm),
```

```
penalty = c("lasso"),
```

```
family = c("gaussian"),
```

```
ncores = numcores, ...
```

# Batch script for R biglasso, highlights

.....

```
#SBATCH --partition=compute
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=64
#SBATCH -A account
... etc
```

```
module load slurm
module load cpu
module load gcc/9.2.0
module load r
```

```
cp Xmatrix_100000_50000.csv /scratch/$USER/job_${SLURM_JOB_ID}/Xinput.csv
cp Ymatrix_100000_1.csv /scratch/$USER/job_${SLURM_JOB_ID}/Yinput.csv
```

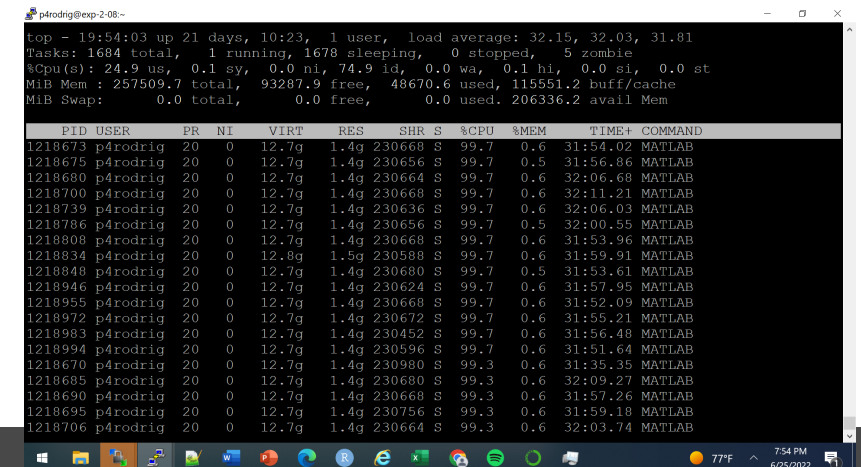
```
Rscript --vanilla path-2-your-script/rLassov2.R
    /scratch/$USER/job_${SLURM_JOB_ID}/Xinput.csv
    /scratch/$USER/job_${SLURM_JOB_ID}/Yinput.csv 64 TRUE > save_std_output.txt
```

# Matlab

- Matlab – has tall table functions to read in large CSV file; but lasso needs tall array

I recalled a 'table2array' function that helped

But matlab kept failing with worker getting disconnected after 1 or 2 hours of processing; might be a memory issue



```
p4rodrig@exp-2-08:~$ top - 19:54:03 up 21 days, 10:23, 1 user, load average: 32.15, 32.03, 31.81
Tasks: 1684 total, 1 running, 1678 sleeping, 0 stopped, 5 zombie
%Cpu(s): 24.9 us, 0.1 sy, 0.0 ni, 74.9 id, 0.0 wa, 0.1 hi, 0.0 si, 0.0 st
MiB Mem : 257509.7 total, 93287.9 free, 48670.6 used, 115551.2 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used, 206336.2 avail Mem

  PID USER      PR  NI   VIRT   RES   SHR  S  %CPU  %MEM     TIME+ COMMAND
1218673 p4rodrig  20   0  12.7g   1.4g  230668 S   99.7   0.6   31:54.02 MATLAB
1218675 p4rodrig  20   0  12.7g   1.4g  230656 S   99.7   0.5   31:56.86 MATLAB
1218680 p4rodrig  20   0  12.7g   1.4g  230664 S   99.7   0.6   32:06.68 MATLAB
1218700 p4rodrig  20   0  12.7g   1.4g  230668 S   99.7   0.6   32:11.21 MATLAB
1218739 p4rodrig  20   0  12.7g   1.4g  230636 S   99.7   0.6   32:06.03 MATLAB
1218786 p4rodrig  20   0  12.7g   1.4g  230656 S   99.7   0.5   32:00.55 MATLAB
1218808 p4rodrig  20   0  12.7g   1.4g  230668 S   99.7   0.6   31:53.96 MATLAB
1218834 p4rodrig  20   0  12.8g   1.5g  230588 S   99.7   0.6   31:59.91 MATLAB
1218848 p4rodrig  20   0  12.7g   1.4g  230680 S   99.7   0.5   31:53.61 MATLAB
1218946 p4rodrig  20   0  12.7g   1.4g  230624 S   99.7   0.6   31:57.95 MATLAB
1218955 p4rodrig  20   0  12.7g   1.4g  230668 S   99.7   0.6   31:52.09 MATLAB
1218972 p4rodrig  20   0  12.7g   1.4g  230672 S   99.7   0.6   31:55.21 MATLAB
1218983 p4rodrig  20   0  12.7g   1.4g  230452 S   99.7   0.6   31:56.48 MATLAB
1218994 p4rodrig  20   0  12.7g   1.4g  230596 S   99.7   0.6   31:51.64 MATLAB
1218670 p4rodrig  20   0  12.7g   1.4g  230980 S   99.3   0.6   31:35.35 MATLAB
1218685 p4rodrig  20   0  12.7g   1.4g  230680 S   99.3   0.6   32:09.27 MATLAB
1218690 p4rodrig  20   0  12.7g   1.4g  230668 S   99.3   0.6   31:57.26 MATLAB
1218695 p4rodrig  20   0  12.7g   1.4g  230756 S   99.3   0.6   31:59.18 MATLAB
1218706 p4rodrig  20   0  12.7g   1.4g  230664 S   99.3   0.6   32:03.74 MATLAB
```

# Dask-ml

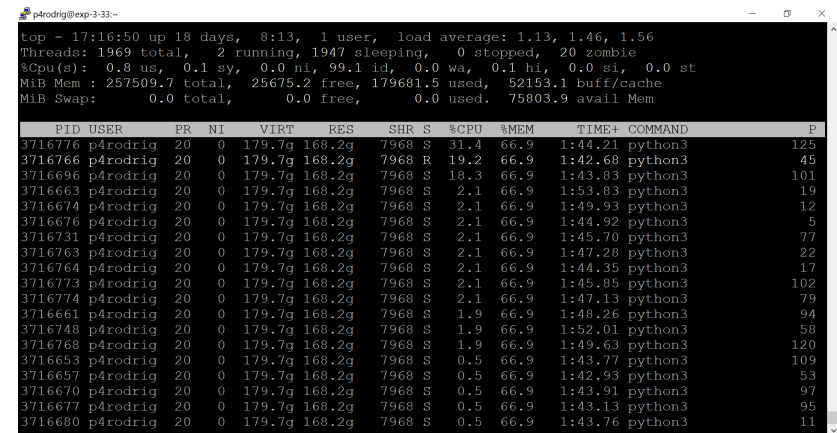
- Dask-ML – worked easily with distributed dataframe, built on top of sklearn,

`read_csv` uses a 'blocksize' parameter

repeatedly calls sklearn 'partial-fit' functions that iterate 1 time for 'batch' processing

worked better with Y&X in one dataframe

I stopped it after 4hrs, but it was running



The image shows a terminal window with the following content:

```
p4rodrig@exp-3-33:~$ top - 17:16:50 up 18 days, 8:13, 1 user, load average: 1.13, 1.46, 1.56
Threads: 1969 total, 2 running, 1947 sleeping, 0 stopped, 20 zombie
%Cpu(s): 0.8 us, 0.1 sy, 0.0 ni, 99.1 id, 0.0 wa, 0.1 hi, 0.0 si, 0.0 st
MiB Mem : 257509.7 total, 25675.2 free, 179681.5 used, 52153.1 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used, 75803.9 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND	P
3716776	p4rodrig	20	0	179.7g	168.2g	7968	S	31.4	66.9	1:44.21	python3	125
3716766	p4rodrig	20	0	179.7g	168.2g	7968	R	19.2	66.9	1:42.68	python3	45
3716696	p4rodrig	20	0	179.7g	168.2g	7968	S	18.3	66.9	1:43.83	python3	101
3716663	p4rodrig	20	0	179.7g	168.2g	7968	S	2.1	66.9	1:53.83	python3	19
3716674	p4rodrig	20	0	179.7g	168.2g	7968	S	2.1	66.9	1:49.93	python3	12
3716676	p4rodrig	20	0	179.7g	168.2g	7968	S	2.1	66.9	1:44.92	python3	5
3716731	p4rodrig	20	0	179.7g	168.2g	7968	S	2.1	66.9	1:45.70	python3	77
3716763	p4rodrig	20	0	179.7g	168.2g	7968	S	2.1	66.9	1:47.28	python3	22
3716764	p4rodrig	20	0	179.7g	168.2g	7968	S	2.1	66.9	1:44.35	python3	17
3716773	p4rodrig	20	0	179.7g	168.2g	7968	S	2.1	66.9	1:45.85	python3	102
3716774	p4rodrig	20	0	179.7g	168.2g	7968	S	2.1	66.9	1:47.13	python3	79
3716661	p4rodrig	20	0	179.7g	168.2g	7968	S	1.9	66.9	1:48.26	python3	94
3716748	p4rodrig	20	0	179.7g	168.2g	7968	S	1.9	66.9	1:52.01	python3	58
3716768	p4rodrig	20	0	179.7g	168.2g	7968	S	1.9	66.9	1:49.63	python3	120
3716653	p4rodrig	20	0	179.7g	168.2g	7968	S	0.5	66.9	1:43.77	python3	109
3716657	p4rodrig	20	0	179.7g	168.2g	7968	S	0.5	66.9	1:42.93	python3	53
3716670	p4rodrig	20	0	179.7g	168.2g	7968	S	0.5	66.9	1:43.91	python3	97
3716677	p4rodrig	20	0	179.7g	168.2g	7968	S	0.5	66.9	1:43.13	python3	95
3716680	p4rodrig	20	0	179.7g	168.2g	7968	S	0.5	66.9	1:43.76	python3	11

# Spark

- Spark – uses distributed dataframe with pyspark.ml.regression package

Use the vector assembler options to get Y and 'features' all in one dataframe

Got some help on spark session options:

device & executor memory, eg: local[64], 8Gb, 2Gb

I stopped it after 4 hrs; it was running

*Local[64], 50 min in, 1 iteration*

```
p4rodrig@exp-6-40:~$ top - 11:34:21 up 18 days, 18:18, 1 user, load average: 1.12, 4.39, 26.17
Threads: 1980 total, 2 running, 1971 sleeping, 0 stopped, 7 zombie
%Cpu(s): 1.0 us, 0.6 sy, 0.0 ni, 98.4 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 257509.7 total, 96845.3 free, 24911.5 used, 135752.9 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used, 230237.8 avail Mem

  PID USER      PR  NI  VIRT  RES  SHR  S  %CPU  %MEM    TIME+  COMMAND
3307477 p4rodrig  20   0  37.3g 15.6g 31100 R  98.3   6.2 15:45.96 Thread-4
3307364 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:55.26 java
3307366 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:55.05 java
3307369 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:55.03 java
3307373 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:55.20 java
3307374 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:55.22 java
3307375 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:55.20 java
3307380 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:55.13 java
3307381 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:55.17 java
3307395 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:54.79 java
3307396 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:54.90 java
3307402 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:54.82 java
3307409 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:55.14 java
3307411 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:54.93 java
3307420 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:54.92 java
3307426 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:54.86 java
3307430 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:54.90 java
3307433 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:55.25 java
3307435 p4rodrig  20   0  37.3g 15.6g 31100 S   1.7   6.2 30:55.20 java
```

# Keras

- Keras - set up a tensorflow dataset with data generator (to reading batches at a time from csv file); a little tricky to get shapes right;

I set up a 1 layer linear neural network with a L1 regularizer on weights

runs, 1000 epochs in 1 hour+ (on cpu)

I did not set up convergence options like other Lasso implementations

# Some notes

- **Start with small data with interactive session or notebook – maybe even just use a smaller sample?**
- **All packages generally work as documented, but**
- **All packages require working through some implementation issues or environment options for the session/job/execution**

# How to use R directly on Expanse

1. Get an interactive compute node:

2. Try

`$ module spider r`      *(this tells you what modules you need)*

3. Enter

`$ module load cpu/0.15.4`

`$ module load gcc/9.2.0`

`$ module load r/4.0.2-openblas`

`$ R`

R version 4.0.2 (2020-06-22) -- "Taking Off Again"

Copyright (C) 2020 The R Foundation for Statistical Computing

Platform: x86\_64-pc-linux-gnu (64-bit)

.....

Type 'q()' to quit R.

>

```
[p4rodrig@login02 ~]$ module spider r
```

```
-----  
r: r/4.0.2-openblas  
-----
```

```
Other possible modules matches:
```

```
AMDuProf, amber, aria2, arm-forge, berkeley-db, bism
```

```
You will need to load all module(s) on any one of the li  
"r/4.0.2-openblas" module is available to load.
```

```
cpu/0.15.4 gcc/9.2.0
```

```
Help:
```



# A note on installing R Packages (into your own directories)

- In R (might help to be on interactive node):

*install.packages('package-name')*

(see <https://cran.r-project.org/> for package lists and reviews)

- Sometimes you have to be explicit:

*install.packages('ggmap',  
 repos='http://cran.us.r-project.org',dependencies=TRUE)*

If compiling is required and you get an error, call support

Packages are put into your /home/user/R directory

## Other R package possibilities:

- Rspark - R interface to Spark
- R Keras – R interface to Keras
- pdbR - distributed matrix support (better for dense matrices vs Spark)
- Rgputools – GPU support

*THE END*